

# A Generalized Risk Budgeting Approach to Portfolio Construction

Martin Haugh <sup>\*</sup> Garud Iyengar <sup>†</sup> Irene Song <sup>‡</sup>

September 2015

## Abstract

Risk-based asset allocation models have received considerable attention in recent years. This increased popularity is due in part to the difficulty in estimating expected returns as well as the financial crisis of 2008 which has helped reinforce the key role of risk in asset allocation. In this study, we propose a generalized risk budgeting (*GRB*) approach to portfolio construction. In a *GRB* portfolio assets are grouped into possibly overlapping subsets and each subset is allocated a pre-specified risk budget. Minimum variance, risk parity and risk budgeting portfolios are all special instances of a *GRB* portfolio. The *GRB* portfolio optimization problem is to find a *GRB* portfolio with an optimal risk-return profile where risk is measured using any positively homogeneous risk measure. When the subsets form a partition, the assets all have the same expected return and we restrict ourselves to long-only portfolios, then the *GRB* problem can in fact be solved as a convex optimization problem. In general, however, the *GRB* problem is a constrained non-convex problem, for which we propose two solution approaches. The first approach uses a semidefinite programming (SDP) relaxation to obtain an (upper) bound on the optimal objective function value. In the second approach we develop a numerical algorithm that integrates augmented Lagrangian and Markov chain Monte Carlo (MCMC) methods in order to find a point in the vicinity of a very good local optimum. This point is then supplied to a standard non-linear optimization routine with the goal of finding this local optimum. It should be emphasized that the merit of this second approach is in its generic nature: in particular, it provides a starting-point strategy for any non-linear optimization algorithm.

**Keywords:** Risk Parity, Risk Budgeting, MCMC, Augmented Lagrangian, Portfolio Optimization, Semidefinite Programming

---

<sup>\*</sup>mh2078@columbia.edu. Department of Industrial Engineering & Operations Research, Columbia University, S.W. Mudd Building, 500 West 120th Street, New York, NY 10027.

<sup>†</sup>garud@ieor.columbia.edu. Department of Industrial Engineering & Operations Research, Columbia University, S.W. Mudd Building, 500 West 120th Street, New York, NY 10027.

<sup>‡</sup>is2306@columbia.edu. Department of Industrial Engineering & Operations Research, Columbia University, S.W. Mudd Building, 500 West 120th Street, New York, NY 10027.

# 1 Introduction

Risk-based asset allocation models have received considerable attention in recent years. Some of this attention has been motivated by the difficulty in estimating expected returns. Mean-variance optimization, for example, is very sensitive to expected asset returns and if applied naively, generally results in portfolios with extreme portfolio weights that are unstable over time. While there are now many methods for addressing these problems, e.g. Black and Litterman (1992), there has been a trend of late to focus on approaches that are more robust to any assumptions on expected returns. The “ $1/N$ ” approach of DeMiguel, Garlappi, and Uppal (2009) is notable in this regard as are the recent developments in risk-based asset allocation models which are the focus of this paper.

As the term “risk-based” suggests, risk generally plays a more important role in risk-based portfolio construction models. Examples of these models include the classic minimum variance approach of Markowitz and the more contemporary risk parity and risk budgeting approaches. In this study, we propose a generalized risk budgeting (*GRB*) approach to portfolio construction.

The concept of risk parity goes back to 1996 when Bridgewater Associates launched a risk parity fund called the All Weather fund. Although the risk parity product was originally introduced to market by Bridgewater, the term “risk parity” was first coined by Qian (2006) who formalized the definition of risk parity in terms of a risk budget where weights of assets are determined in such a way that they all contribute equally to the overall portfolio risk. Maillard, Roncalli, and Teiletche (2010) referred to such a portfolio as an equal risk contribution (ERC) portfolio. They analyzed properties of an unconstrained long-only ERC portfolio and showed that its volatility lies between the volatilities of the long-only minimum variance and equally-weighted portfolios. We note here that the terms “risk parity” and “equal risk contribution” are used interchangeably in the literature, but hereafter we will use the former.

A risk parity portfolio, however, is not always desirable. An investor may prefer to allocate different risk budgets to each asset, and this preference would require a more general risk budgeting portfolio. Theoretical properties of risk budgeting portfolios were analyzed by Bruder and Roncalli (2012). Extending the result of Maillard et al. (2010), they showed that the volatility of a long-only risk budgeting portfolio lies between the volatilities of a long-only minimum variance portfolio and a long-only weighted portfolio whose weights are proportional to their risk budgets. They further demonstrated that when the portfolio risk is computed using a convex risk measure and risk budgets are defined to be strictly positive, a long-only risk budgeting portfolio exists and is unique.

Since the introduction of this approach, there have been many additional studies on risk parity and risk budgeting approaches. Most of them, however, have focused on seeking a long-only minimal risk portfolio that satisfies (pre-defined) risk budgeting constraints. The majority of these methods, therefore, lack flexibility. For example, by disregarding the expected asset returns in their problem formulations, many of these methods make the implicit assumption that all asset returns are identical in expectation. Whether or not the disregarding of expected returns results in a better-performing portfolio, such an assumption does not hold in practice. In addition, it is clearly desirable that investors be able to freely express their views on expected asset returns when constructing a portfolio.<sup>1</sup>

In this paper we propose a generalized risk budgeting (*GRB*) problem formulation that leads

<sup>1</sup>Research that addresses the issue of incorporating expected asset returns into risk constrained portfolio selection includes Zhu, Li, and Sun (2010); Zhu, X. T. Gui, Sun, and Li (2011/2012); Boudt, Carl, and Peterson (2012). Zhu et al. (2010) studied optimal mean-variance marginal risk constrained portfolio selection with the purpose of controlling relative risk contribution of individual assets; Zhu et al. (2011/2012) presented a factor-risk-constrained mean-variance portfolio optimization model; and, Boudt et al. (2012) proposed a minimum conditional-value-at-risk (CVaR) concentration portfolio that balances the investor’s return objective and the diversification of risk across portfolio constituents.

in general to a non-convex optimization problem. We refer to this problem as the *GRB* portfolio optimization problem. We then develop solution approaches for this *GRB* problem.<sup>2</sup> The key advantage of our formulation over the prevailing risk parity or risk budgeting approaches is that it offers a much greater degree of flexibility in the way risk-based portfolios are constructed. It allows for short sales of assets, the use of risk factors to model asset returns, and most importantly, it allows investors to define risk budgets for overlapping subsets of assets.

When the subsets form a partition, the assets all have the same expected return and we restrict ourselves to long-only portfolios, we show that the problem can be formulated as a convex optimization problem and is therefore easily solved. This result generalizes Bruder and Roncalli (2012)’s approach for constructing a long-only risk budgeting portfolio with minimum variance. For the more general *GRB* problem, we propose two solution approaches. The first approach is a semidefinite programming (SDP) relaxation to the problem that allows us to obtain an (upper) bound on the optimal objective function of the *GRB* problem. We remark that the solution to this SDP problem also often yields a very good starting point for a generic non-linear optimization solver. To our knowledge, we are not aware of any other studies that apply an SDP relaxation to the risk parity or risk budgeting problems. Our second approach develops a numerical algorithm that combines augmented Lagrangian and Markov chain Monte-Carlo (MCMC) methods with the goal of finding a point in the vicinity of a very good local optimum. This point is then supplied to a non-linear optimization routine to compute this local optimum. The merit of this second approach is in its generic nature: in particular, it provides a starting-point strategy for any non-linear optimization algorithm.

The remainder of the paper is organized as follows. In Section 2 we formally define the *GRB* problem and describe our two solution approaches as well as the special case that can be solved as a convex optimization problem. We provide numerical and empirical results for the SDP relaxation and the augmented Lagrangian-MCMC approach in Section 3, and discuss some of the practical challenges associated with applying the latter approach in Section 4. We then conclude in Section 5.

## 2 The Generalized Risk Budgeting (*GRB*) Problem

In portfolio construction and analysis it is often preferable to group assets according to attributes such as asset class, country, sector and industry. In the case of an investment portfolio with a broad coverage of asset classes, for example, it may be more insightful and therefore preferable to look at the marginal risk contribution of each asset class rather than each individual asset in the portfolio. Similarly, in the case of a large-scale stand-alone asset class portfolio, it may be preferable to control the risk contributions of assets at an aggregate level such as country or market sector. The generalized risk budgeting (*GRB*) strategy is based on this very idea of managing the marginal risk contributions of subsets of assets to the total portfolio risk. In a *GRB* portfolio, the risk contribution from each (pre-specified) subset of assets is set equal to some pre-specified risk budget. Note that we are using the term “subset” rather than “partition” since depending on the attributes used for the asset classification, assets may belong to more than one group. We will see later that minimum variance, risk parity and risk budgeting portfolios are all special instances of a *GRB* portfolio.

The objective of the *GRB* problem is to find a *GRB* portfolio that is optimal on the basis of its risk-return profile. Portfolio risk in the *GRB* problem is computed via a positively homogeneous risk measure for which we can use Euler’s theorem to provide a risk decomposition. Examples of positively homogeneous risk measures include portfolio volatility, value-at-risk (VaR) and any

---

<sup>2</sup>Henceforth, we refer to the *GRB* portfolio optimization problem as the *GRB* problem.

coherent risk measures such as conditional value-at-risk (CVaR) (Artzner, Delbaen, Eber, and Heath, 1999).

Towards this end, let  $\mathcal{R}(x) : \mathbb{R}^d \rightarrow \mathbb{R}$  denote a generic risk measure that is a positively homogeneous function of degree one in the portfolio weight vector,  $x$ . Euler's theorem then provides the following additive risk decomposition:

$$\mathcal{R}(x) = \sum_{i=1}^d x_i \frac{\partial \mathcal{R}(x)}{\partial x_i} \quad (2.1)$$

where the marginal risk contribution of the  $i$ -th asset is

$$RC_i(x) = x_i \frac{\partial \mathcal{R}(x)}{\partial x_i}.$$

If  $\mathcal{M}_1, \dots, \mathcal{M}_s \subseteq \{1, \dots, d\}$  denote  $s$  subsets of portfolio assets,<sup>3</sup> then the marginal risk contribution of the  $k$ -th subset is

$$RC_{\mathcal{M}_k}(x) := \sum_{i \in \mathcal{M}_k} RC_i(x).$$

Let  $\beta_1, \dots, \beta_s$  now denote the risk budgets for  $\mathcal{M}_1, \dots, \mathcal{M}_s$ , respectively. We can then formulate the *GRB* problem as:

$$\begin{aligned} (GRB) \quad & \max_{x \in \mathcal{X}} \quad \mu'x - \lambda \mathcal{R}(x), \\ & \text{subject to} \quad \sum_{i \in \mathcal{M}_k} RC_i(x) = \beta_k \mathcal{R}(x), \quad k = 1, \dots, s \end{aligned}$$

where  $\mu \in \mathbb{R}^d$  is a vector of expected returns,  $\lambda$  is a risk aversion parameter and  $\mathcal{X} := \{x \in \mathbb{R}^d : \mathbf{1}'x = 1\}$ . Note that the constraint  $\sum_{i \in \mathcal{M}_k} RC_i(x) = \beta_k \mathcal{R}(x)$  implies that  $\sum_{k=1}^s \beta_k = 1$  when the  $\mathcal{M}_k$ 's form a partition. We note that the *GRB* problem becomes a minimum variance problem when  $\mu = \mu_0 \mathbf{1}$ , there is only one subset  $\mathcal{M}_1$  which is equal to the universe of assets, and the risk measure is portfolio volatility. It is a risk parity problem when  $\mu = \mu_0 \mathbf{1}$ , the  $\mathcal{M}_k$ 's are all singletons and all risk budgets,  $\beta_k$ , are equal. Finally it is a risk budgeting problem when  $\mu = \mu_0 \mathbf{1}$  and the  $\mathcal{M}_k$ 's are again all singletons.

The *GRB* problem is a constrained non-convex optimization problem for which efficient solution algorithms are unavailable. Although there are numerous methods available for computing risk parity portfolio weights, e.g. Spinu (2013) and Bai, Scheinberg, and Tutuncu (2013), these methods are in general not applicable to the *GRB* problem. In Section 2.1 below we consider a special case of the *GRB* problem which can be solved as a convex optimization problem. We then proceed to discuss our solution approaches for the general non-convex case. Note that the parameter and variable notation introduced in this section will be used throughout the paper unless otherwise stated.

---

<sup>3</sup>In practice, the subsets  $\mathcal{M}_k$  would typically correspond to the different asset classes included in an investment portfolio. For example, AQR Risk Parity funds contain approximately 90 assets from 6 different asset classes – fixed income, TIPS, equity, currency, commodity and credit. The size of each asset class in these funds ranges from 5 to 25 securities. For the complete holdings of AQR Risk Parity funds, refer to <http://funds.aqr.com/our-funds/global-allocation-funds/risk-parity-fund>.

## 2.1 A Special Case of the GRB Problem

We now consider a special case of the GRB problem in which all assets have the same expected return, i.e.  $\mu = \mu_0 \mathbf{1}$ , each asset belongs to one and only one subset, i.e. the  $\mathcal{M}_k$ 's form a partition of the asset space, and non-negativity constraints are imposed at the partition level. We can then reformulate the GRB problem as follows:

$$\begin{aligned} & \min_{x \in \mathcal{X}} \mathcal{R}(x), \\ & \text{subject to } \sum_{i \in \mathcal{M}_k} RC_i(x) = \beta_k \mathcal{R}(x), \quad k = 1, \dots, s, \\ & \sum_{i \in \mathcal{M}_k} x_i \geq 0, \quad k = 1, \dots, s. \end{aligned} \quad (2.2)$$

Assuming also that each  $\beta_k > 0$ , we then have the following result that extends Bruder and Roncalli (2012).

**Theorem 2.1.** *Assuming  $\mathcal{R}(y) \neq 0$  for nonzero  $y$ , then problem (2.2) is equivalent to the convex optimization problem (2.3):*

$$\begin{aligned} & \min_y \mathcal{R}(y) \\ & \text{subject to } \sum_{k=1}^s \beta_k \ln \left( \sum_{i \in \mathcal{M}_k} y_i \right) \geq c \end{aligned} \quad (2.3)$$

where  $c$  is an arbitrary constant. In particular the normalized optimal solution  $\tilde{y}^*$  to (2.3) is also the optimal solution to (2.2). (A normalized solution is one where  $\sum_{i=1}^d \tilde{y}_i^* = 1$ . See the discussion after the proof.)

*Proof.* Let  $\mathcal{L}(y, \gamma)$  denote the Lagrangian of the optimization problem (2.3) so that

$$\mathcal{L}(y, \gamma) = \mathcal{R}(y) - \gamma \left( \sum_{k=1}^s \beta_k \ln \left( \sum_{i \in \mathcal{M}_k} y_i \right) - c \right).$$

At optimality, the solution  $y^*$  satisfies the KKT conditions. That is,  $y^*$  satisfies (i) the first-order conditions

$$\frac{\partial \mathcal{L}(y, \gamma)}{\partial y_i} = \frac{\partial \mathcal{R}(y)}{\partial y_i} - \gamma \left( \frac{\beta_k}{\sum_{j \in \mathcal{M}_k} y_j} \right) = 0, \quad (2.4)$$

for  $i = 1, \dots, d$  and where  $k$  is the index of the subset  $\mathcal{M}_k$  containing  $i$ , and (ii) the complementary slackness conditions

$$\gamma \left( \sum_{k=1}^s \beta_k \ln \left( \sum_{i \in \mathcal{M}_k} y_i \right) - c \right) = 0. \quad (2.5)$$

Note that as  $\ln : \mathbb{R}^+ \rightarrow \mathbb{R}$ , we must have  $\sum_{i \in \mathcal{M}_k} y_i > 0$  for  $k = 1, \dots, s$ , and hence  $y$  cannot be  $\mathbf{0}$ . Then since  $\mathcal{R}(y) \neq 0$  for nonzero  $y$ , at least one of  $\frac{\partial \mathcal{R}(y)}{\partial y_i}$  must be nonzero by (2.1). The strict positivity of the  $\beta_k$ 's and (2.4) then imply  $\gamma > 0$ . We therefore have

$$\frac{\partial \mathcal{R}(y)}{\partial y_i} = \gamma \left( \frac{\beta_k}{\sum_{j \in \mathcal{M}_k} y_j} \right) \quad (2.6)$$

for  $i = 1, \dots, d$  and where  $\gamma > 0$ . Multiplying both sides of (2.6) by  $y_i$  and then summing over

$i \in \mathcal{M}_k$  yields

$$\begin{aligned} \sum_{i \in \mathcal{M}_k} y_i \frac{\partial \mathcal{R}(y)}{\partial y_i} &= \gamma \left( \frac{\beta_k}{\sum_{j \in \mathcal{M}_k} y_j} \right) \sum_{i \in \mathcal{M}_k} y_i \\ &= \gamma \beta_k \end{aligned}$$

for  $k = 1, \dots, s$ . We therefore see that the risk contribution of each  $\mathcal{M}_k$  is proportional to its risk budget,  $\beta_k$ . The normalized optimal solution  $\tilde{y}^*$  is then the optimal solution  $x^*$  to (2.2) as claimed.  $\blacksquare$

Note that as  $\sum_{k=1}^s \beta_k \ln \left( \sum_{i \in \mathcal{M}_k} y_i \right) = c$  by (2.5), we could directly obtain  $x^*$  from (2.3) if we used  $c^* = c - \ln \left( \sum_{i=1}^d y_i \right)$  in (2.3), rather than the original  $c$  which led to the solution  $y$ . Note also that we recover the results of Bruder and Roncalli (2012) if the  $\mathcal{M}_k$  are all singletons.

## 2.2 An SDP Relaxation for the General GRB Problem

Our first approach for the *GRB* problem uses a semidefinite programming (SDP) relaxation to obtain an upper bound on the optimal objective function value. There are two advantages of the SDP approach: (i) the solution to the SDP problem (which is generally infeasible for the *GRB* problem) can be used as a (hopefully very good) starting point for a standard non-linear optimization routine, and (ii) the SDP solution can often provide a “certificate” of near-optimality when the SDP solution has an objective function that is close to the objective function of the best local optimal solution that we have found.

In our development of the SDP approach we will assume initially that our risk measure is portfolio volatility so that

$$\mathcal{R}(x) := \sqrt{x' \Sigma x}$$

where  $\Sigma \in \mathbb{R}^{d \times d}$  is the covariance matrix of asset returns. The marginal risk contributions of the individual assets then satisfy

$$RC_i(x) = x_i \frac{(\Sigma x)_i}{\sqrt{x' \Sigma x}}, \quad i = 1, \dots, d.$$

With this measure of risk we can rewrite the *GRB* problem in the following equivalent form:

$$\begin{aligned} &\max_{x, X} \quad \mu'x - \lambda \mathcal{R}(x), \\ \text{subject to} \quad &\sum_{i \in \mathcal{M}_k} \text{tr}(\Gamma_i X) = \beta_k \text{tr}(\Sigma X), \quad k = 1, \dots, s, \\ &X = xx', \\ &\mathbf{1}'x = 1 \end{aligned} \tag{2.7}$$

where  $\Gamma_i = e_i e_i' \Sigma$ ,  $e_i$  denotes the  $i$ -th column of the identity matrix  $I \in \mathbb{R}^{d \times d}$ ,  $x \in \mathbb{R}^d$ , and  $\text{tr}(\cdot)$  denotes the trace of a matrix. Since  $X = xx'$  is the only non-convex constraint in (2.7), we obtain a convex relaxation of the *GRB* problem by relaxing this constraint to  $X \succeq xx'$ . We then obtain

the following SDP relaxation of our *GRB* problem:

$$\begin{aligned}
& \max_{x, X} \mu'x - \lambda \mathcal{R}(x), \\
& \text{subject to } \sum_{i \in \mathcal{M}_k} \text{tr}(\Gamma_i X) = \beta_k \text{tr}(\Sigma X), \quad k = 1, \dots, s \\
& \quad \begin{bmatrix} X & x \\ x' & 1 \end{bmatrix} \succeq 0, \\
& \quad \mathbf{1}'x = 1
\end{aligned} \tag{2.8}$$

where we have used Schur complement to reformulate the semidefinite constraint  $X \succeq xx'$  as a linear matrix inequality. Note that we can recover (2.7) from (2.8) by imposing an additional (non-convex) constraint that the left-hand-side of the LMI in (2.8) to be a rank one matrix. The SDP relaxation can be solved efficiently and the SDP solution provides an upper bound on the optimal objective function of the *GRB* problem. For example, one can easily implement and solve (2.8) using<sup>4</sup> **CVX** (Grant and Boyd, 2014, 2008).

SDP relaxations can also be formulated for the *GRB* problem with other risk measures. Suppose the risk measure is either the value-at-risk (VaR) or conditional value-at-risk (CVaR) of a portfolio. Let  $F_r(z) := P\{r \leq z\}$  denote the CDF of the portfolio return  $r$ , then the VaR and CVaR at the confidence level  $\alpha \in (0, 1)$  are defined as

$$VaR_\alpha := \min\{z | F_r(z) \geq \alpha\}$$

and

$$CVaR_\alpha := E[r | r \geq VaR_\alpha(r)].$$

Suppose the asset returns are, as before, normally distributed with mean vector  $\mu$  and covariance matrix  $\Sigma$ . Then for a portfolio with the weight vector  $x$  it easily follows that

$$VaR_\alpha(x) = \mu'x + \Psi_1(\alpha)\sqrt{x'\Sigma x}$$

and

$$CVaR_\alpha(x) = \mu'x + \Psi_2(\alpha)\sqrt{x'\Sigma x}$$

where (see McNeil, Frey, and Embrechts (2005))

$$\begin{aligned}
\Psi_1(\alpha) &= \sqrt{2} \text{erf}^{-1}(2\alpha - 1), \\
\Psi_2(\alpha) &= (\sqrt{2\pi} \exp(\text{erf}^{-1}(2\alpha - 1))^2 (1 - \alpha))^{-1}, \\
\text{erf}(z) &= (2/\sqrt{\pi}) \int_0^z e^{-t^2} dt.
\end{aligned}$$

Likewise, the respective marginal VaR and CVaR contributions<sup>5</sup> of the  $i$ -th asset are given by

$$RC_{i, VaR_\alpha}(x) = \mu_i + x_i \frac{(\Sigma x)_i}{\sqrt{x'\Sigma x}} \Psi_1(\alpha)$$

<sup>4</sup>A special SDP mode in **CVX** allows positive (negative) semidefinite constraints  $\succeq$  ( $\preceq$ ) to be imposed using **Matlab**'s standard inequality operators  $>=$  ( $<=$ ).

<sup>5</sup>See Boudt et al. (2012), who also studied portfolio selection under CVaR budgets.

and

$$RC_{i,CVaR_\alpha}(x) = \mu_i + x_i \frac{(\Sigma x)_i}{\sqrt{x' \Sigma x}} \Psi_2(\alpha).$$

Without loss of generality, let us consider the risk budgeting constraints that arise when we use VaR as our risk measure. These constraints take the form

$$\sum_{i \in \mathcal{M}_k} \left( \mu_i + x_i \frac{(\Sigma x)_i}{\sqrt{x' \Sigma x}} \Psi_1(\alpha) \right) = \beta_k \left( \mu' x + \Psi_1(\alpha) \sqrt{x' \Sigma x} \right), \quad k = 1, \dots, s. \quad (2.9)$$

The VaR version of the *GRB* problem is given by:

$$\begin{aligned} & \max_{x, X} \mu' x - \lambda VaR_\alpha(x) \\ & \text{subject to } \sum_{i \in \mathcal{M}_k} \left( \mu_i \sqrt{x' \Sigma x} + x_i (\Sigma x)_i \Psi_1(\alpha) \right) = \beta_k \left( \mu' x \sqrt{x' \Sigma x} + \Psi_1(\alpha) x' \Sigma x \right), \quad k = 1, \dots, s \\ & \quad X = x x' \\ & \quad \mathbf{1}' x = 1. \end{aligned} \quad (2.10)$$

Note that the risk budgeting constraints of (2.10) are obtained by multiplying both sides of (2.9) by  $\sqrt{x' \Sigma x}$ . Let  $t = \sqrt{x' \Sigma x}$ . Then the optimization problem (2.10) can be reformulated as follows:

$$\begin{aligned} & \max_{t, \delta, w, x, y, X, Z} \mu' x - \lambda VaR_\alpha(x) \\ & \text{subject to } \sum_{i \in \mathcal{M}_k} (\mu_i t + \text{tr}(\Gamma_i X) \Psi_1(\alpha)) = \beta_k (\mu' \delta + \Psi_1(\alpha) \text{tr}(\Sigma X)), \quad k = 1, \dots, s, \\ & \quad Z = \begin{bmatrix} 1 & t & x' \\ t & y & \delta' \\ x & \delta & X \end{bmatrix}, \quad \text{rank}(Z) = 1, \\ & \quad t = \sqrt{x' \Sigma x}, \\ & \quad t, y \geq 0, x, \delta \in \mathbb{R}^d, X \in \mathbb{R}^{d \times d}, Z \succeq 0, \end{aligned} \quad (2.11)$$

where we use the fact that  $\text{rank}(Z) = 1$  implies that  $\delta = tx$ , and  $X = xx'$ . There are two non-convex constraints in (2.11):  $\text{rank}(Z) = 1$  and  $t = \sqrt{x' \Sigma x}$ . One can obtain a semidefinite relaxation for the feasible set by dropping the  $\text{rank}(Z) = 1$  constraint, and relaxing the other constraint to  $t \geq \sqrt{x' \Sigma x}$ <sup>6</sup>. Since  $VaR_\alpha(x)$  is not a convex function of  $x$ , one will have to replace  $VaR_\alpha(x)$  by a concave function that is pointwise larger in order to obtain a valid relaxation for (2.11). This last step would not be necessary if we used  $CVaR$  as our risk measure because  $CVaR(x)$  is already a convex function of  $x$ . We also note that while we have assumed normally distributed asset returns here, other distributions such as the  $t$  distribution could be used instead. See Boyd and Vandenberghe (1997) for SDP relaxations of non-convex problems, more generally.

### 2.3 An Augmented Lagrangian-MCMC Approach

Our second approach to solving the *GRB* problem involves combining the augmented Lagrangian approach with MCMC sampling to generate a point in the proximity of the *global* optimum of the *GRB* problem. This point can then be used as a starting point for a non-linear optimization routine to converge to a globally optimal *GRB* portfolio. The underlying idea of the algorithm is to effectively

<sup>6</sup>This is a second-order cone constraint that can in turn be reformulated as a semidefinite constraint.



sample points with a higher objective function value and simultaneously drive the sample path in the direction of the feasible region using the augmented Lagrangian terms.

Let  $\Omega$  be the state space and  $p(x) = C^{-1}p^*(x)$  denote some target probability distribution on  $\Omega$  where  $C := \int_{\Omega} p^*(x)dx$  is the normalization constant. The MCMC method is an approach to sample from  $p(x)$  when the normalizing constant is hard compute. In the MCMC approach, one constructs a Markov chain on  $\Omega$  using a “proposal” distribution  $q(x_{t+1}|x_t)$  in such a way that  $p(x)$  is the unique stationary distribution for the Markov chain. Modulo some technical conditions<sup>7</sup>, the main requirement of MCMC is that the unnormalized distribution,  $p^*(x)$ , should be easy to compute. Given a current sample  $x_t$  at time  $t$  the proposal distribution,  $q(\cdot|x_t)$ , is used to generate a candidate sample,  $x_{t+1}$ , which is then accepted with probability

$$\alpha(x_t, x_{t+1}) := \min \left\{ 1, \frac{q(x_t|x_{t+1})p^*(x_{t+1})}{q(x_{t+1}|x_t)p^*(x_t)} \right\}. \quad (2.12)$$

If the candidate point  $x_{t+1}$  is rejected we then set  $x_{t+1} = x_t$  and continue sampling in this manner.

Since our goal is to solve the *GRB* problem, one possibility would be to set

$$p^*(x) = \exp(\gamma F(x)) \mathbb{I}_{\mathcal{F}}(x)$$

where  $F(x) := \mu'x - \lambda \mathcal{R}(x)$  denotes the objective function of the *GRB* problem,  $\gamma$  is an annealing parameter that is used to concentrate the  $p^*(x)$  in the proximity of the global optimum, and  $\mathbb{I}_{\mathcal{F}}(\cdot)$  denotes the indicator function of the set  $\mathcal{F} = \{x \in \mathcal{X} \mid h_k(x) := \sum_{i \in \mathcal{M}_k} RC_i(x) - \beta_k \mathcal{R}(x) = 0, k = 1, \dots, s\}$ . Since the feasible region  $\mathcal{F}$  of the *GRB* problem is, typically, very “small”;  $p^*(x_{t+1})$  is likely to be zero, for most candidate points  $x_{t+1}$  and these points will be rejected in the acceptance-rejection step (2.12). Therefore, using MCMC to sample *only* from the feasible region is very difficult, and particularly so for high-dimensional problems.

One possible approach to overcoming these difficulties is to allow the MCMC iterates  $x_t$  to be *infeasible*, but to “direct” them towards the feasible region by adding a term which penalizes infeasibility to our definition of  $p^*$ . In particular, we could define

$$P_c(x) := F(x) + \frac{1}{2}c\|h(x)\|_2^2$$

where  $c$  is a negative constant and now use  $p^*(x) = e^{\gamma P_c(x)}$  as the unnormalized density. The main difficulty with the penalty approach is that it is very sensitive to the value of the penalty parameter  $c$ . This is a well-known phenomenon, and the augmented Lagrangian algorithm was introduced in order to circumvent this numerical instability.

<sup>7</sup>See Robert and Casella (2004) for further technical details on MCMC algorithms.

### The AL-MCMC Algorithm

In the augmented Lagrangian approach, we define the time<sup>8</sup>  $t$  target distribution to be  $p_t^*(x) := e^{\gamma_t \mathcal{L}_{c_t}(u_t, x)}$  where the augmented Lagrangian function of the *GRB* problem is defined as:

$$\begin{aligned} \mathcal{L}_{c_t}(u_t, x) &:= F(x) + u_t' h(x) + \frac{1}{2} c_t \|h(x)\|_2^2 \\ &= \mu' x - \lambda \mathcal{R}(x) + \sum_{k=1}^s u_{t,k} \left( \beta_k \mathcal{R}(x) - \sum_{i \in \mathcal{M}_k} RC_i(x) \right) \\ &\quad + \frac{1}{2} c_t \left( \sum_{k=1}^s \left( \beta_k \mathcal{R}(x) - \sum_{i \in \mathcal{M}_k} RC_i(x) \right)^2 \right) \end{aligned}$$

where  $u_t = (u_{t,1}, \dots, u_{t,s}) \in \mathbb{R}^s$  is a vector of time  $t$  Lagrange<sup>9</sup> multipliers. Let  $d_{c_t}(u) := \max_{x \in \mathcal{X}} \mathcal{L}_{c_t}(u, x)$  denote the dual objective.

The initial vector of Lagrange dual multipliers  $u_0$  and the penalty parameter  $c_0$  are specified exogenously. The values for dual multipliers  $u_t$  and the non-increasing penalty parameter  $c_t$  for  $t \geq 1$  are chosen adaptively during the course of the simulation. In particular, we decrease  $c_t$  by a predetermined value  $\epsilon_c$  when there is no improvement in constraint violations over a particular iteration. When there is an improvement in constraint violation, we do not update  $c_t$  but instead update the Lagrange multipliers using the first order conditions, i.e. we set

$$u_{t+1} = u_t - \epsilon_u \nabla d_{c_t}(u_t) \quad (2.13)$$

where  $d_{c_t}(u)$  denotes the dual function, and  $\epsilon_u$  is a given step size. We chose not to update both  $c_t$  and  $u_t$  in every iteration in order to ensure that we leave the current location only after adequately exploring its neighborhood. The update of duals  $u_t$  or the penalty parameter  $c_t$  occurs every iteration whether or not the candidate  $x_{t+1}$  is accepted. We note that (2.13) represents the steepest descent iteration for minimizing  $d_{c_t}$  but one may choose other methods such as Newton's method for updating the Lagrange multipliers (see Appendix A). Furthermore, one can use other criteria for updating  $c_t$ . See Bertsekas (1996) for a detailed discussion of the augmented Lagrangian method.

For generating a candidate value of  $x_{t+1}$ , we use a proposal distribution based on a random walk chain. In particular, we generate  $z_t^* \sim N(0, \sigma_t^q I)$  and take

$$x_{t+1} = x_t + z_t^*$$

as our candidate point which is then accepted with probability  $\alpha(x_t, x_{t+1})$ . The standard deviation of the proposal distribution  $\sigma_t^q$  can be thought of as a tuning parameter that we can adjust to increase the acceptance probability when the current location is near a feasible region. This is done by decreasing  $\sigma_t^q$  by a factor of  $\kappa$ , where  $0 < \kappa < 1$ , only if the percentage drop in the size of the constraint violations is larger than a predetermined value  $\delta$ .<sup>10</sup> The value of  $\kappa$  thus can be

<sup>8</sup>We note that since  $p_t^*$  now changes with each iteration, there is no longer a fixed target stationary distribution for our algorithm.

<sup>9</sup>See Appendix A for further details on the augmented Lagrangian functions.

<sup>10</sup>See Algorithm 1 for precise details. Depending on the specific problem under consideration, one may choose to modify this step or simply to keep  $\sigma_t^q$  constant across all  $t$ . However, considering that choosing the acceptance rate is important for a good numerical performance of the algorithm, it is recommended that one allows  $\sigma_t^q$  to vary with  $t$ . Also, instead of having the annealing schedule, one can directly adjust  $\sigma_t^q$  based on the acceptance rate at time  $t$ . This is in fact considered as common practice for adjusting the mixing and the acceptance probability of the chain.

interpreted as the rate at which we decrease  $\sigma_t^q$  when the current location is considered favorable. If  $\kappa$  is too small, then  $\sigma_t^q$  decreases too fast and moves become small too quickly. Small moves are generally accepted (high acceptance probability), and as a result, the chain is prone to get trapped in the current region prior to exploring other regions. On the contrary, if  $\kappa$  is too big, then  $\sigma_t^q$  decreases too slowly and moves remain large. This leads to many rejections resulting in an inefficient chain.<sup>11</sup>

In each iteration, irrespective of whether the proposed sample  $x_{t+1}$  is accepted or rejected, the annealing parameter  $\gamma_t$  is increased according to:

$$\gamma_t = \sigma_\gamma \gamma_{t-1} \quad (2.14)$$

where  $\sigma_\gamma = \left(\frac{\gamma_{max}}{\gamma_0}\right)^{\frac{1}{T}}$  and  $\gamma_{max}$  is the maximum allowed value of  $\gamma_t$ . Note that (2.14) is known as the geometric annealing schedule.<sup>12</sup> Thus, the AL-MCMC algorithm is a simulated annealing algorithm<sup>13</sup> where by forcing  $\lim_{t \rightarrow \infty} \gamma_t = \infty$  we hope to drive samples towards the global optimum of the *GRB* problem.

The AL-MCMC algorithm attempts to combine the best aspects of the augmented Lagrangian method and the MCMC method. The augmented Lagrangian term guides the Markov chain towards a feasible region, while the acceptance-rejection step in the MCMC method attempts to ensure that the iterates do not get trapped in poor local maxima of the *GRB* problem.

A complete specification of our AL-MCMC algorithm is given in Algorithm 1. A feasible sample point with the highest value of  $F(\cdot)$  is probably most suitable to be used as a starting point for a non-linear optimization routine. However, as the direct sampling of a feasible point is overly difficult for the *GRB* problem, the last point obtained by the algorithm is then fed to a non-linear optimization routine with the goal of quickly finding a good nearby local maximum.

We note that our algorithm is a heuristic algorithm that we hope is capable of producing good starting points for a non-linear optimization solver. We expect this algorithm to be useful for general non-convex optimization problems beyond the *GRB* problem of this paper. There is also further scope for improvement. For example, we could use a more sophisticated MCMC algorithm as compared to the Metropolis-Hastings. For example, if we suspect that  $F(\cdot)$  or  $L_{ct}(\cdot)$  is multimodal then hybrid MCMC methods such as Hamiltonian MCMC should be superior. It is also possible to tailor the proposal distributions,  $q(x_{t+1}|x_t)$ , for the problem at hand. Note also that while it is not explicitly stated, it of course makes sense to keep track of the best feasible sample that has been obtained during the execution of the algorithm.

We therefore propose the following procedure to solve the *GRB* problem:

**Step 1.** Generate an initial vector  $x_0$  to be used as the starting point of the Markov chain and choose values of  $\gamma_0, \sigma_\gamma, \epsilon_c, \epsilon_u, \delta, \sigma_0^q, c_0, u_0$  and  $\kappa$  to be used as parameters for the AL-MCMC algorithm (Algorithm 1).

---

The basic idea is to increase  $\sigma_t^q$  when the acceptance rate is too high and decrease  $\sigma_t^q$  when the acceptance rate is too low. Refer to Roberts, Gelman, and Gilks (1994) for further details.

<sup>11</sup>In choosing  $\kappa$ , we generated a training data set for each test case presented in Section 3, and tested multiples of 0.25 as candidate values of  $\kappa$ . Based on this, we selected  $\kappa = .75$  across all test cases.

<sup>12</sup>The geometric annealing schedule is the most commonly used annealing schedule in practice, also suggested by the originators of the simulated annealing algorithm (Kirkpatrick, Gelatt, and Vecchi, 1983). Depending on the specific problem under consideration, one may choose to use a different annealing schedule. For a comparison of different annealing schedules, refer to Nourani and Andresen (1998).

<sup>13</sup>See Van Laarhoven and Arts (1987) for further details.

**Algorithm 1** AL-MCMC

---

```

1: Choose  $x_0, \gamma_0, \sigma_\gamma, \epsilon_c, \epsilon_u, \delta, \sigma_0^q, c_0, u_0, \kappa$ .
2: for  $t = 0 : T$  do
3:   Draw a candidate sample  $x_{t+1}$  from the proposal  $q(x_{t+1}|x_t)$ .
4:   Let  $\alpha(x_t, x_{t+1}) = \min \left\{ 1, \frac{q(x_t|x_{t+1})p^*(x_{t+1})}{q(x_{t+1}|x_t)p^*(x_t)} \right\}$  where  $\frac{p^*(x_{t+1})}{p^*(x_t)} = e^{\gamma_t(\mathcal{L}_{c_t}(u_t, x_{t+1}) - \mathcal{L}_{c_t}(u_t, x_t))}$ .
5:   if  $\alpha \geq 1$  then
6:      $x_{t+1} \leftarrow x_{t+1}$  # Accept the candidate
7:   else
8:     Draw  $p \sim \mathbb{U}[0, 1]$ 
9:     if  $p \leq \alpha$  then
10:       $x_{t+1} \leftarrow x_{t+1}$  # Accept the candidate
11:    else
12:       $x_{t+1} \leftarrow x_t$  # Reject the candidate
13:    end if
14:  end if
15:   $\gamma_{t+1} \leftarrow \sigma_\gamma \gamma_t$  # Update the annealing parameter
16:  if  $\|h(x_{t+1})\|_2^2 < \|h(x_t)\|_2^2$  then
17:     $u_{t+1} \leftarrow u_t - \epsilon_u \nabla d_{c_t}(u_t)$  where  $d_{c_t}(u) = \max_{x \in \mathcal{X}} \mathcal{L}_{c_t}(u, x)$  # Update the Lagrange multipliers
18:    if  $\frac{\|h(x_t)\|_2^2}{\|h(x_{t+1})\|_2^2} - 1 > \delta$  then
19:       $\sigma_{t+1}^q \leftarrow \kappa \sigma_t^q$  # Update the jump size
20:    end if
21:  else
22:     $c_{t+1} \leftarrow c_t + \epsilon_c$  # Update the penalty parameter
23:  end if
24: end for

```

---

**Step 2.** Perform the AL-MCMC algorithm to obtain an initial point  $x_s$  to be fed to a non-linear optimization routine.

**Step 3.** Solve the *GRB* problem using a non-linear optimization solver with  $x_s$  obtained from Step 1 as the initial guess.

The AL-MCMC algorithm described in this section can be further enhanced by using a set of different random starting points  $x_0$  for generating Markov chains. For instance, in our numerical experiments we used antithetic starting points to generate several values of  $x_s$ <sup>14</sup>.

### 3 Numerical Results

We now present numerical results for the two proposed approaches: the SDP relaxation and the AL-MCMC algorithm. We first describe a simple example with the goal of illustrating the potential effectiveness of the AL-MCMC algorithm. We then turn to discuss the performances of the two approaches when they are tested on *GRB* problems with the number of assets ranging from 7 to 200. All the results presented in this section are based on percentage returns, i.e. returns are multiplied by 100, unless otherwise stated. Note also that the term “optimal solution” generally

<sup>14</sup>Readers interested in antithetic variates in Monte Carlo techniques can refer to Robert and Casella (2004).

denotes a local optimum.

### 3.1 Numerical Results for a Small Example

Our first problem<sup>15</sup> is a 5-asset problem with a variance-covariance matrix of percentage returns:

$$\Sigma = \begin{bmatrix} 94.868 & 33.750 & 12.325 & -1.178 & 8.778 \\ 33.750 & 445.642 & 98.955 & -7.901 & 84.954 \\ 12.325 & 98.955 & 117.265 & 0.503 & 45.184 \\ -1.178 & -7.901 & 0.503 & 5.460 & 1.057 \\ 8.778 & 84.954 & 45.184 & 1.057 & 34.126 \end{bmatrix}.$$

We also assumed that the expected returns of these assets are identical so that  $\mu = \mu_0 \mathbf{1}$ . Suppose now we want to compute a long-only risk parity portfolio with minimum variance and that we apply the AL-MCMC algorithm to solve this problem. We used a single Markov chain of 5,000 points, i.e.  $T = 5,000$  in Algorithm 1, and  $x_0$  was generated uniformly from the 5-dimensional unit cube. We also used the following parameters:

- initial annealing parameter  $\gamma_0 = 1$  with  $\sigma_\gamma = 1.0007$ ;
- initial penalty parameter  $c_0 = -10,000$  with  $\epsilon_c = 0$ ;
- jump size  $\sigma_0^q = 0.5$  with  $\kappa = 0.75$ ;
- threshold parameter for updating  $\sigma_t^q$ ,  $\delta = 0.01$ ;
- initial Lagrange multipliers  $u_0 = \mathbf{0}$  with  $\epsilon_u = 0.01$ .

Since this problem is relatively simple with just five constraints, we did not need to update  $c$  during the course of the algorithm.  $x_{5000} = [0.1245; 0.0467; 0.0833; 0.6133; 0.1323]$  is the last point obtained from the AL-MCMC algorithm. If we specify the feasibility tolerance to  $10^{-4}$ , this point is, in fact, the optimal risk parity solution.<sup>16</sup> Without the use of a non-linear optimization routine, the AL-MCMC algorithm was therefore able to discover a good risk parity solution. The running time<sup>17</sup> for the algorithm was 1.31 seconds.

When the algorithm was applied without the penalty parameter, i.e.  $c_t = 0$  for all  $t$ , or without the Lagrange multipliers, i.e.  $u_t = \mathbf{0}$  for all  $t$ , it failed to converge to a risk parity solution. All of its sampled points violated the risk parity constraints by more than  $10^{-4}$ , and hence, the help of a non-linear optimization routine was necessary for finding an optimal risk parity solution. When its last point was supplied to a non-linear optimization routine, the optimal risk parity solution was found successfully. These results demonstrate the potential advantage of incorporating the augmented Lagrangian method into the MCMC algorithm.

### 3.2 Numerical Results for the GRB Problem

For more general GRB problems we focused on the portfolio volatility risk measure  $\mathcal{R}(x) := \sqrt{x' \Sigma x}$  and assumed a risk aversion parameter  $\lambda$  of 1. Expected asset returns  $\mu$ , covariance matrices  $\Sigma$ , and risk budgets  $\beta$  are all generated randomly. In particular, we sampled  $\mu$  from  $N(0, I)$ , and for  $\Sigma$ ,

<sup>15</sup>This is the same example presented in Bai et al. (2013).

<sup>16</sup>One can readily check that  $x^*$  is indeed the optimal risk parity solution by solving Problem 2.3 directly.

<sup>17</sup>All our experiments were performed using `Matlab` on an Intel Core i5-680 (3.60GHz), 64-bit operating system.

we first generated a matrix  $V \in \mathbb{R}^{d \times d}$  using a standard normal distribution, and then converted it into a symmetric positive semidefinite matrix by multiplying it by its transpose; i.e.  $\Sigma = V'V$ . We generated risk-budgets  $\beta = (\beta_1, \dots, \beta_s)$  from  $\mathbb{U}^s(0, 1)$  and normalized them such that  $\sum_{k=1}^s \beta_k = 1$ .

We considered the five test cases listed in Table 1 under two different scenarios. In the first scenario, we assumed that  $\mu = \mu_0 \mathbf{1}$ , i.e. all assets have identical returns. In the second scenario, we allowed the assets to have different returns. Also, within each of these five test cases, we considered two test sets. In the first test set (Set 1) the subsets of assets form a partition whereas in the second test set (Set 2) we considered overlapping subsets of assets that did not form a partition. The subsets  $\mathcal{M}_k$  for  $k = 1, \dots, s$ , for each test set were chosen randomly by generating an  $s \times d$  matrix  $M$  of zeros and ones. A nonzero entry, say  $M_{(i,j)}$ , then indicates that the  $j$ -th asset belongs to the  $i$ -th subset ( $\mathcal{M}_i$ ). In the case of the first test set, nonzero entries that result in overlapping subsets of assets were simply set to zero so that the resulting subsets formed a partition.

Table 1: Test Case Descriptions

Test Case	Number of assets ( $d$ )	Number of subsets ( $s$ )	Max. function evals. for <code>fmincon</code>
1	7	3	3,000
2	30	5	5,000
3	50	5	7,000
4	100	10	20,000
5	200	10	60,000

In order to evaluate the AL-MCMC algorithm, we generated 5 antithetic pairs of random points. For each pair  $(x, x')$  of random points, we first sampled  $x = (x_1, \dots, x_d)$  from  $\mathbb{U}^d(0, 1)$  and set  $x' = (1 - x_1, \dots, 1 - x_d)$ . We used each of these 10 points as the starting point  $x_0$  to generate  $T = 1,000$  samples from the Markov chain. Therefore, we simulated a total of  $N = 10$  Markov chains.<sup>18</sup> The final<sup>19</sup> point, i.e. the  $T$ -th point, from each chain is then used as the initial starting point of a non-linear optimization solver. In our experiments we used the `fmincon` solver with the interior point method in `Matlab` as our non-linear optimization solver. The maximum number of function evaluations allowed for `fmincon` for each test case is specified in the final column of Table 1. We also used `Matlab` and `CVX` (Grant and Boyd, 2014, 2008) for solving the SDP relaxation of the *GRB* problem (2.8).

### 3.2.1 Identical Returns: $\mu = \mu_0 \mathbf{1}$

In this scenario, we assume  $\mu = \mu_0 \mathbf{1}$ . In this case, the *GRB* problem reduces to the minimum risk problem subject to the risk budgeting constraints.

We solved the *GRB* problem using four different approaches. First, we solved the SDP relaxation to obtain the lower bound. Next, we use the possibly infeasible solution of the SDP relaxation as the initial point for `fmincon`. We refer to this as the SDP-`fmincon` approach. Next, we solved the problem using the AL-MCMC-`fmincon` approach, i.e. we simulated 10 Markov chains starting from 10 random initial points generated using the antithetic random variate method, and used the  $T = 1,000^{th}$  iterate of each chain as the initial point for a call to `fmincon`. In order to benchmark

<sup>18</sup>The values of  $N$  and  $T$  were chosen by testing various combinations on a training data set for each test case. In practice, the number of MCMC iterations typically ranges from 1,000 to 10,000 combined across all chains used; e.g.  $N = 10$  chains with a length of  $T = 1,000$  each corresponds  $N \times T = 10,000$  total iterations. Refer to Section 4 for a discussion on how to choose chain lengths,  $T$ , and the number of chains,  $N$ .

<sup>19</sup>The rationale behind choosing the final point is that due to the risk budgeting constraints, sampled points are most likely to be infeasible, and an infeasible sample point attaining the highest value of  $F(\cdot)$  is not necessarily the best point in terms of its proximity to the optimum solution.

the contribution of the MCMC algorithm, we solved the *GRB* problem using **fmincon** starting from  $10 \times T = 10,000$  random starting points distributed according to  $\mathbb{U}^d(0, 1)$ .<sup>20</sup> In addition, we also considered the alternating linearization backtracking (ALM-BTKR) approach to solve risk parity problems. The ALM-BTKR approach was introduced by Bai et al. (2013) where the risk parity problem was formulated as the quadratic least-squares problem:

$$\begin{aligned} \min_{x \in \mathcal{X}, \theta} \quad & \sum_{k=1}^s (\sum_{i \in \mathcal{M}_k} x_i (\Sigma x)_i - \theta)^2 \\ \text{subject to} \quad & a_i \leq x_i \leq b_i, \quad i = 1, \dots, d. \end{aligned} \quad (3.1)$$

Note that risk parity is achieved when (3.1) has an optimal value of zero. This approach can be easily extended to the case where the risk of the various asset classes  $\mathcal{M}_k$  are not equal by scaling  $\theta$  as follows:

$$\begin{aligned} \min_{x \in \mathcal{X}, \theta} \quad & \sum_{k=1}^s (\sum_{i \in \mathcal{M}_k} x_i (\Sigma x)_i - \beta_k \theta)^2 \\ \text{subject to} \quad & a_i \leq x_i \leq b_i, \quad i = 1, \dots, d. \end{aligned} \quad (3.2)$$

We stress that (3.1) and (3.2) are only able to identify a portfolio that satisfies the risk-parity constraints; they do not seek a risk-parity portfolio with minimum risk.

In Table 2 we report the following metrics for each of the four solution methods: **SDP-fmincon**, **AL-MCMC-fmincon**, **fmincon** and **ALM-BTKR**.

- $\min \tilde{F}(x^*)$  denotes the objective value of the best feasible portfolio computed by the methods. For the **SDP-fmincon** we report the risk of the portfolio computed by **fmincon** starting from the SDP solution, if feasible. For the **AL-MCMC-fmincon** method we report the best objective value among all feasible solutions resulting from the 10 starting points, and for the **fmincon** approach we report the best objective value among all feasible portfolio resulting from the  $10 \times T = 10,000$  random starting points. For **ALM-BTKR**, we report  $\tilde{F}(x_{ALM-BTKR}^*)$  where  $x_{ALM-BTKR}^*$  is the solution obtained by solving (3.2).
- The range of objective values of all the feasible solutions computed by the method. We do not report a solution range for **ALM-BTKR** and **SDP-fmincon** since these methods yield at most one feasible solution.
- The SDP lower bound.
- The number of failures. A failure occurs when **fmincon** does not return a feasible solution for a feasibility tolerance of  $10^{-6}$ .
- $t$  (sec). We report the amount of time taken to obtain  $\min \tilde{F}(x^*)$  for the first time over the 10 trials for the **AL-MCMC-fmincon** algorithm and over the 10,000 trials<sup>21</sup> for the **fmincon** algorithm. For **ALM-BTKR** and **SDP-fmincon**,  $t$  represents the total execution time as each only yields a single solution.

The main issue with reporting the execution time  $t$  in the manner described above is that, in order to determine the  $\min \tilde{F}$ , we first need to compute all 10 solutions for the **AL-MCMC-fmincon** algorithm and 10,000 solutions for the **fmincon** algorithm. However, given that **AL-MCMC-fmincon** uses only 10 random points and **fmincon** uses 10,000 random points, comparing the total execution time of each method is fraught with difficulties. The advantage of reporting  $t$  in the above way is

<sup>20</sup>As solving *GRB* problems with **fmincon** was extremely time-consuming, 1,000 random starting points were used instead of 10,000 for Test Case 5 with 200 assets.

<sup>21</sup>See footnote 20.

that it allows for a fairer comparison of AL-MCMC-`fmincon` with `fmincon`. Note that all solutions are reported with a precision of four decimal places.

From the results reported in Table 2, the portfolio  $x_{ALM-BTKR}^*$  is not the minimum variance portfolio. In most cases,  $\tilde{F}(x_{ALM-BTKR}^*)$  is at least 50% larger than the solutions obtained by the other two methods. Also, due to the backtracking component of ALM-BTKR, the method may become prohibitively slow for higher dimensional problems. For Test Case 1 with 7 assets the ALM-BTKR method takes less than 25 seconds, but for Test Case 4 with 100 assets the execution time is more than 350 seconds. In fact, for Test Case 5 with 200 assets, ALM-BTKR fails to find a feasible solution even after 2,350 seconds.



Table 2: Numerical Results for the Case of  $\mu = \mu_0 \mathbf{1}$ 

In this table, we report results for four methods: SDP relaxation, SDP-fmincon, AL-MCMC-fmincon, fmincon, and ALM-BTKR for the case where  $\mu = \mu_0 \mathbf{1}$ . The first and second columns contain the test set number and the name of the algorithm, respectively. The third column reports the best solution obtained. For the ALM-BTKR method, this column reports  $\tilde{F}(x_{ALM-BTKR}^*)$  where  $x_{ALM-BTKR}^*$  is the solution obtained by solving (3.2) using the ALM-BTKR method. The fourth column reports the range of the obtained solutions. The fifth column reports the lower bound on the objective function  $\tilde{F}(x)$  obtained by the SDP relaxation. The sixth column reports the number of failures. The last column reports the execution time in seconds required to obtain the best solution.

Test Case 1: 7 Assets and 3 Subsets						
Set	Method	$\min \tilde{F}(x^*)$	Solution range	SDP lower bound	No. of failures	$t$ (sec)
1	SDP-fmincon	33.17	–	31.72	0	3.64
1	AL-MCMC-fmincon	33.09	[33.09, 34.16]	31.72	0	18.50
1	fmincon	33.09	[33.09, 229.82]	31.72	194	100.26
1	ALM-BTKR	111.00	–	31.72	0	16.84
2	SDP-fmincon	25.76	–	25.57	0	2.06
2	AL-MCMC-fmincon	25.78	[25.78, 31.91]	25.57	0	16.90
2	fmincon	25.77	[25.77, 98.65]	25.57	840	1,764.15
2	ALM-BTKR	N/A	–	25.57	1	21.37

Test Case 2: 30 Assets and 5 Subsets						
Set	Method	$\min \tilde{F}(x^*)$	Solution range	SDP lower bound	No. of failures	$t$ (sec)
1	SDP-fmincon	38.59	–	38.06	0	3.44
1	AL-MCMC-fmincon	38.59	[38.59, 38.59]	38.06	0	2.54
1	fmincon	38.59	[38.59, 328.61]	38.06	968	5.48
1	ALM-BTKR	100.61	–	38.06	0	49.39
2	SDP-fmincon	40.12	–	39.55	0	1.26
2	AL-MCMC-fmincon	40.12	[40.12, 40.12]	39.55	0	0.81
2	fmincon	40.12	[40.12, 40.12]	39.55	0	0.50
2	ALM-BTKR	N/A	–	39.55	1	63.71

Test Case 3: 50 Assets and 5 Subsets						
Set	Method	$\min \tilde{F}(x^*)$	Solution range	SDP lower bound	No. of failures	$t$ (sec)
1	SDP-fmincon	38.24	–	36.97	0	3.65
1	AL-MCMC-fmincon	38.24	[38.24, 38.24]	36.97	0	2.68
1	fmincon	38.25	[38.25, 328.33]	36.97	2,339	4.17
1	ALM-BTKR	79.89	–	36.97	0	106.72
2	SDP-fmincon	64.33	–	51.77	0	2.71
2	AL-MCMC-fmincon	64.33	[64.33, 64.33]	51.77	0	1.90
2	fmincon	64.33	[64.33, 64.33]	51.77	0	1.41
2	ALM-BTKR	N/A	–	51.77	1	118.22

Test Case 4: 100 Assets and 10 Subsets						
Set	Method	$\min \tilde{F}(x^*)$	Solution range	SDP lower bound	No. of failures	$t$ (sec)
1	SDP-fmincon	55.33	–	53.76	0	14.84
1	AL-MCMC-fmincon	55.33	[55.33, 103.99]	53.76	0	7.98
1	fmincon	55.33	[55.33, 314.89]	53.76	3,213	26.88
1	ALM-BTKR	110.39	–	53.76	0	354.29
2	SDP-fmincon	51.71	–	48.28	0	14.02
2	AL-MCMC-fmincon	51.71	[51.71, 51.71]	48.28	0	7.51
2	fmincon	51.71	[51.71, 51.71]	48.28	0	6.22
2	ALM-BTKR	N/A	–	48.28	1	423.02

Test Case 5: 200 Assets and 10 Subsets						
Set	Method	$\min \tilde{F}(x^*)$	Solution range	SDP lower bound	No. of failures	$t$ (sec)
1	SDP-fmincon	54.48	–	49.85	0	127.43
1	AL-MCMC-fmincon	54.48	[54.48, 54.68]	49.85	5	67.89
1	fmincon	54.48	[54.48, 393.48]	49.85	642	613.94
1	ALM-BTKR	N/A	–	49.85	1	2,607.56
2	SDP-fmincon	56.02	–	49.85	0	127.69
2	AL-MCMC-fmincon	55.02	[55.02, 55.23]	49.85	0	54.87
2	fmincon	55.02	[55.02, 110.96]	49.85	58	114.84
2	ALM-BTKR	N/A	–	49.85	1	2,378.06

When all assets are assumed to have identical expected returns, the SDP relaxation appears to provide a fairly effective lower bound, against which we can compare solutions obtained from other methods. For example, in Set 2 of Test Case 1, the difference between the SDP relaxation and the AL-MCMC-`fmincon` solution is just 0.21, and therefore, we know that the AL-MCMC solution is close to the global optimum. Since the *GRB* problem is non-convex, having an effective lower bound on its objective function is very informative. The SDP-`fmincon` method moreover exhibits comparable performance to the AL-MCMC-`fmincon` method in most cases except for Set 1 of Test Case 1, in which the optimal solution of the SDP-`fmincon` method is 0.08 higher than that of the AL-MCMC-`fmincon` method.

The AL-MCMC-`fmincon` method generally has better performance than `fmincon`. The execution time of `fmincon` is inconsistent. For example, the execution time of `fmincon` for Test Case 1 with 7 assets ranges anywhere from 100 seconds to 1,764 seconds to obtain the best solution. For Set 1 of Test Case 3 and Test Case 4, `fmincon` failed over 2,000 and 3,000 times, respectively. Similarly, for Set 1 Test Case 5, it failed over 64% (642/1,000) of the time. In contrast, the AL-MCMC-`fmincon` method was able to find an optimal solution within 20 seconds for Test Case 1 – 4, and in the majority of cases, it took less than 10 seconds to do so. For Test Case 5, we observe significant increase in execution time across all four methods. However, AL-MCMC-`fmincon`'s increased execution time of about a minute is still superior to the other methods. In addition, for Test Case 1 – 4, the AL-MCMC-`fmincon` method *never* failed thereby suggesting that the AL-MCMC approach is able to generate a good initial point for `fmincon`. The magnitude of constraint violations that caused its failures in Set 1 Test Case 5 was, moreover, less than  $10^{-4}$ . The range of values of the solutions computed by `fmincon` is also much wider than the range for AL-MCMC-`fmincon`. The test cases in which subsets form a partitions, i.e. Set 1, clearly highlight the fact that `fmincon` significantly underperforms AL-MCMC-`fmincon`. For instance, in Set 1 Test Case 4, the `fmincon` solution range is [55.33, 314.89] whereas the AL-MCMC-`fmincon` solutions range is [55.33, 103.99]. This is actually the widest solution range we see for AL-MCMC-`fmincon`. In other test cases, solution ranges are very narrow for the AL-MCMC-`fmincon` algorithm.

We also compared the performance of the AL-MCMC algorithm against AL-MCMC but with the Lagrange multipliers  $u_t \equiv 0$  for all  $t$ , i.e. a pure penalty method, and AL-MCMC with the penalty parameter  $c_t \equiv 0$  for all  $t$ , i.e. a pure dual method. We found that the AL-MCMC-`fmincon` algorithm overall performed better than both of these alternatives. These results further demonstrate the merit of integrating the augmented Lagrangian method with the MCMC algorithm. We report the results for the AL-MCMC algorithm with  $u_t \equiv 0$  and  $c_t \equiv 0$  in Appendix B.

### 3.2.2 General Expected Returns

The next set of results are for the case where expected returns are not identical across assets. Table 3 presents upper bounds for the *GRB* problem that were obtained using the SDP relaxation. We also report the maximum constraint violation

$$\max_{k=1,\dots,s} \left| \sum_{i \in \mathcal{M}_k} \frac{RC_i(x_{SDP}^*)}{\mathcal{R}(x_{SDP}^*)} - \beta_k \right| \quad (3.3)$$

of the optimal SDP solution,  $x_{SDP}^*$ .

Unlike the previous section, we see that the maximum constraint violation of the solution of the SDP relaxation is quite large in certain cases. For example, it is 0.76 for Set 1 Test Case 1, 0.63 for Set 2 Test Case 2 and 0.52 for Set 1 Test Case 3<sup>22</sup>. In these cases, the SDP upper bound

<sup>22</sup>Recall that the  $\beta_k$ 's are all positive and sum to 1 so that a violation of 0.76 is indeed quite large.

Table 3: Numerical Results for the SDP Relaxation

Table 3 presents the upper bound on the objective function,  $F(x) = \mu'x - \mathcal{R}(x)$ , of the *GRB* problem obtained via the SDP relaxation. The first and second columns contain the test set and case number, respectively. The third and fourth columns contain the number of assets ( $d$ ) and the number of the subsets ( $s$ ) considered in each test case, respectively. The sixth column contains the maximum constraint violation of the optimal SDP solution, i.e. (3.3), and the final column reports the execution time (in seconds) of *CVX* for solving the SDP relaxation.

Set	Test Case	$d$	$s$	Upper bound	Max. constraint violation	$t$ (sec.)
1	1	7	3	103.82	0.76	0.38
2	1	7	3	33.68	0.39	0.38
1	2	30	5	-22.73	0.24	0.58
2	2	30	5	28.38	0.63	0.98
1	3	50	5	5.41	0.52	0.92
2	3	50	5	-3.51	0.39	0.87
1	4	100	10	-15.62	0.10	6.29
2	4	100	10	3.25	0.23	6.66
1	5	200	10	24.05	0.27	87.38
2	5	200	10	24.05	0.35	70.21

is likely to be slack since the SDP solutions are far from being feasible.

In Table 4 we report the following metrics for **SDP-fmincon**, **AL-MCMC-fmincon**, and **fmincon** starting from 10,000 random points.

- $\max F(x^*)$  denotes the objective value of the best feasible solution computed by the three solution methods.
- The range of values for all feasible solutions. Note that we do not report a solution range for **SDP-fmincon** since this method produces only one solution.
- The SDP upper bound.
- A failure occurs when **fmincon** does not return a feasible solution at a feasibility tolerance of  $10^{-6}$ .
- $t$  (sec). We report the amount of time taken to obtain  $\max F(x^*)$  for the first time in the 10 trials for the **AL-MCMC-fmincon** algorithm and the 10,000 trials<sup>23</sup> for the **fmincon** algorithm.

All solutions are reported with a precision of four decimal places. It is clear that the overall performance of the **AL-MCMC-fmincon** method is superior to the other two methods. In comparison to the **SDP-fmincon** approach, we note that the **AL-MCMC-fmincon** method was able to find an optimal solution to all test cases. The **SDP-fmincon** method failed to find even a feasible solution for Set 1 Test Case 1 and Set 1 Test Case 3. It is interesting to note that these failures occurred when the maximum constraint violations of the SDP relaxation solutions were noticeably large (see Table 3). This suggests that an upper bound obtained from the SDP relaxation may turn out to be slack when the optimal SDP solution violates the risk budgeting constraints by a large amount. In the previous section, we saw that the SDP relaxation provides a relatively tight bound when all assets have identical expected returns.

In comparison to **fmincon**, **AL-MCMC-fmincon** exhibits a more stable and consistent performance. It is apparent from Table 4 that the solution ranges given by **fmincon** are very wide in general. For example, the objective values of the solutions obtained from **fmincon** for Set 1 of

<sup>23</sup>See footnote 20.

Table 4: Numerical Results for the Case of  $\mu \neq \mu_0 \mathbf{1}$ 

Table 4 presents the numerical results for the three methods: **SDP-fmincon** method, **AL-MCMC-fmincon** method and **fmincon**, when  $\mu \neq \mu_0 \mathbf{1}$ . The first and second columns contain the test set number and the name of the algorithm, respectively. The third column reports the best solution, i.e.  $\max F(x^*)$ , and the fourth column reports the range of the obtained solutions. The fifth column reports the upper bound on the objective function  $F(x)$  obtained by the SDP relaxation. The sixth column reports the number of failures. The final column reports the execution time (in seconds).

Test Case 1: 7 Assets and 3 Subsets						
Set	Method	$\max F(x^*)$	Solution range	SDP upper bound	No. of failures	$t$ (sec)
1	SDP-fmincon	N/A	–	103.82	1	2.03
1	AL-MCMC-fmincon	63.89	[49.42, 63.89]	103.82	0	4.48
1	fmincon	63.88	[–235.38, 63.88]	103.82	990	6,398.85
2	SDP-fmincon	30.31	–	33.68	0	0.88
2	AL-MCMC-fmincon	30.31	[30.31, 30.31]	33.68	0	0.58
2	fmincon	30.31	[–83.04, 30.31]	33.68	568	2.17
Test Case 2: 30 Assets and 5 Subsets						
Set	Method	$\max F(x^*)$	Solution range	SDP upper bound	No. of failures	$t$ (sec)
1	SDP-fmincon	–25.33	–	–22.73	0	1.91
1	AL-MCMC-fmincon	–25.33	[–25.33, –25.33]	–22.73	0	2.37
1	fmincon	–25.33	[–309.20, –25.33]	–22.73	640	15.13
2	SDP-fmincon	14.89	[14.89, 14.89]	28.37	0	1.50
2	AL-MCMC-fmincon	14.89	[14.89, 14.89]	28.37	0	0.68
2	fmincon	14.89	[–14.05, 14.89]	28.37	0	0.42
Test Case 3: 50 Assets and 5 Subsets						
Set	Method	$\max F(x^*)$	Solution range	SDP upper bound	No. of failures	$t$ (sec)
1	SDP-fmincon	N/A	–	5.41	1	N/A
1	AL-MCMC-fmincon	–0.30	[–0.30, –0.30]	5.41	0	2.52
1	fmincon	–0.30	[–345.38, –0.30]	5.41	2,253	2.30
2	SDP-fmincon	–20.72	[–20.72, –20.72]	28.38	0	2.42
2	AL-MCMC-fmincon	–20.72	[–20.72, –20.72]	28.38	0	2.00
2	fmincon	–20.72	[–20.72, –20.72]	28.38	0	1.35
Test Case 4: 100 Assets and 10 Subsets						
Set	Method	$\max F(x^*)$	Solution range	SDP upper bound	No. of failures	$t$ (sec)
1	SDP-fmincon	–17.28	–	–15.62	0	13.97
1	AL-MCMC-fmincon	–17.27	[–17.28, –17.27]	–15.62	0	7.43
1	fmincon	–17.27	[–302.76, –17.27]	–15.62	1,878	8.43
2	SDP-fmincon	–0.46	[–0.46, –0.46]	28.38	0	13.26
2	AL-MCMC-fmincon	–0.46	[–0.46, –0.46]	28.38	0	7.61
2	fmincon	–0.46	[–43.29, –0.46]	28.38	0	6.87
Test Case 5: 200 Assets and 10 Subsets						
Set	Method	$\max F(x^*)$	Solution range	SDP upper bound	No. of failures	$t$ (sec)
1	SDP-fmincon	8.41	–	24.05	0	135.16
1	AL-MCMC-fmincon	8.41	[–172.35, 8.41]	24.05	4	47.55
1	fmincon	8.42	[–379.73, 8.42]	24.05	535	735.36
2	SDP-fmincon	7.50	–	24.05	0	106.25
2	AL-MCMC-fmincon	7.52	[7.51, 7.52]	24.05	0	40.66
2	fmincon	7.52	[–98.34, 7.52]	24.05	6	296.05

Test Case 1 range from –235.38 to 63.89, and it took 6,398.85 seconds to discover the best of these solutions despite Test Case 1 being a low dimensional problem with only 7 assets. This suggests that when **fmincon** uses an unfavorable starting point the solution it obtains can be very far from a good local optimum. The results of Set 1 Test Case 3 further demonstrate this. Of the 10,000 random starting points, **fmincon** failed to find a feasible solution 2,253 times. In many cases, the stand-alone **fmincon** has a relatively large number of failures compared to **AL-MCMC-fmincon**, and hence, one needs to try a very large number of random starting points. In contrast,

the AL-MCMC-`fmincon` method yields solutions whose ranges are much narrower. Except for Set 1 of Test Case 1 and Test Case 5, all the obtained solutions are very close to the best solution. Moreover, except for Set 1 Test Case 5, the AL-MCMC-`fmincon` method never fails to produce a feasible point. This means that the starting points generated by the AL-MCMC algorithm are much more favorable than random starting points. We remark that even those failures in Set 1 Test Case 5 were caused by the constraint violations that are less than  $10^{-4}$ . Also note that in most cases, AL-MCMC-`fmincon` finds an optimal solution in less than 10 seconds. For Test Case 5, AL-MCMC-`fmincon` solves the *GRB* problem within 50 seconds, which is quite exceptional compared to the other two methods. Based on these observations we can conclude that AL-MCMC-`fmincon` appears to be a much more reliable tool for solving the general *GRB* problem with non-identical expected asset returns.

### 3.3 Empirical Results for the *GRB* Problem

As noted in Section 2, it is often advisable when evaluating the risk of a large-scale equity portfolio to group stocks according to attributes, such as market sector. Accordingly, risk-based sector weighting is a good example of a risk budgeting strategy that can be built using *GRB* models in practice. We now present some empirical results for such an application.

We took our security universe to be the 200 largest (by market capitalization) stocks in the S&P 500. We assigned these stocks to ten sectors – Consumer Discretionary, Consumer Staples, Energy, Financials, Healthcare, Industrials, Information Technology, Materials, Telecommunications Services, and Utilities – according to the Global Industry Classification Standard (GICS) and estimated the monthly covariance matrix<sup>24</sup> over the past 10 years (520 weeks) from August 5<sup>th</sup>, 2005 to July 17<sup>th</sup>, 2015. Our goal in each month was to compute a minimum variance portfolio that satisfies the pre-defined risk-budgeting constraints at a market sector level.<sup>25</sup> We, therefore, have a total of 120 empirical test cases. We note that these test cases are equivalent to Set 1 Test Case 5 presented in Section 3.2.1. We also note that due to the large number of test cases, we use 10 instead of 10,000 randomly generated starting points for `fmincon`, and keep all other testing parameters the same as those used in our earlier numerical experiments.

In Table 5, we report the following metrics for each of the three solution methods: SDP-`fmincon`, AL-MCMC-`fmincon`, and `fmincon`.<sup>26</sup>

- The number of successes. The  $i$ -th trial is deemed a success for a given method if the method achieves the minimum objective value  $v_i^*$  on that trial. For SDP-`fmincon` the objective value is defined as the risk of the portfolio computed by `fmincon` starting from the SDP solution, provided the SDP solution is feasible. For AL-MCMC-`fmincon` and `fmincon` the objective value is the defined as the least risk among all feasible solutions resulting from the 10 starting points.
- The mean % deviation (MPD) from the optimal objective value  $v^*$ . The MPD is defined by the formula:

$$MPD = \frac{1}{120} \sum_{i=1}^{120} \frac{1}{|P_i|} \sum_{j \in P_i} \frac{\tilde{F}(x_{i,j}) - v_i^*}{v_i^*}$$

<sup>24</sup>To estimate the covariance matrix in any given month, we used the previous 2 years of weekly returns and fit a one-factor model based on the first principal component of the return data. For further details on the covariance matrix estimation using factor models, see Ruppert (2010).

<sup>25</sup>Risk-budgets  $\beta = (\beta_1, \dots, \beta_{10})$  are generated as  $\mathbb{U}^{10}(0, 1)$  and then normalized so that  $\sum_{k=1}^{10} \beta_k = 1$ .

<sup>26</sup>ALM-BTKR is not included in our empirical testing as its execution time for the 200-asset case exceeded 2,350 seconds, and despite such execution time, it failed to find a feasible portfolio.

where  $P_i$  is the set of feasible portfolios,  $x_{i,j}$ , computed by the method under consideration from the 10 starting points and  $v_i^*$  is the best objective value on the  $i$ -th test case. We do not report this metric for **SDP-fmincon** since this method yields at most one feasible solution.

- The average failure rate (%). A failure occurs when **fmincon** does not return a feasible solution for a feasibility tolerance of  $10^{-6}$ . For **AL-MCMC-fmincon** and **fmincon** we report the average % of times feasible solutions were not found in the 10 trials across 120 test cases. For **SDP-fmincon** we report the average % of times feasible solutions were not found across 120 test cases.
- The average time (sec). For **AL-MCMC-fmincon** and **fmincon** we report the average amount of time taken to obtain the optimal objective value for the *first* time in the 10 trials across 120 test cases. For **SDP-fmincon**,  $t$  represents the average execution time across 120 test cases as each only yields a single solution.

Table 5: Empirical Testing Results

Table 5 presents the empirical testing results for the three methods: **SDP-fmincon** method, **AL-MCMC-fmincon** method and **fmincon**, when  $\mu = \mu_0 \mathbf{1}$ . The first column contains the name of the algorithm. The second column reports the number of successes, and the third column reports the mean % deviation of a feasible solution from the optimal objective value  $v^*$ . The fourth column reports the average failure rate, and the final column reports the average execution time (in seconds).

Method	No. of successes	Mean % dev. from $v^*$	Average failure rate (%)	Average time (sec)
<b>SDP-fmincon</b>	5	–	77.50	109.94
<b>AL-MCMC-fmincon</b>	120	77.16	21.33	137.19
<b>fmincon</b>	96	97.31	49.92	192.21

**AL-MCMC-fmincon** is successful in *all* test cases, i.e. the risk of the solution computed by **AL-MCMC-fmincon** is the lowest for each test case. Moreover, for most test cases, **AL-MCMC-fmincon** found an optimal solution in the shortest time. Although the reported average execution time for **SDP-fmincon** is less than that of **AL-MCMC-fmincon**, it does not account for the fact that **SDP-fmincon** failed to find a feasible solution 93 times out of 120 test cases. The failure rate of **AL-MCMC-fmincon** (21.33%) is significantly lower than that of the other two methods (77.50% and 49.92% for **SDP-fmincon** and **fmincon** respectively). In terms of solution quality, the MPD from  $v^*$  for **AL-MCMC-fmincon** is approximately 20% lower than that for **fmincon**. This shows that **AL-MCMC-fmincon** is better suited for solving the *GRB* as when compared to **fmincon**.

The empirical results presented here are consistent with the numerical results presented in Section 3. The **AL-MCMC-fmincon** method is a more robust and dependable approach for solving the *GRB* problems.

## 4 Discussion

It is well known that non-linear optimization methods perform poorly, and can even fail, if the starting point is poorly chosen. When solving a multimodal problem, in particular, these methods could get easily trapped by a (poor) local optimum due to an unfavorable starting point. The **AL-MCMC** approach proposed in this study combines the augmented Lagrangian, MCMC and simulated annealing to avoid this phenomenon. We expect that the **AL-MCMC** method can be adapted to solve other multimodal problems.

The **AL-MCMC** method has several parameters that need to be specified including, but not limited to, the chain length, the number of chains, the variance  $\sigma_t^q$  of the proposal distribution, and

the annealing schedule. As with any numerical method, the performance of AL-MCMC does depend on the choice of parameters. For example, the selection of  $\sigma_t^q$  can greatly affect the performance of the sampler. The most widely used approach to choosing  $\sigma_t^q$  is the trial-and-error approach. In this approach, initial values of the tuning parameters are chosen and modified based on the examination of the mixing properties of the chains.<sup>27</sup> For instance, although the overall acceptance rate in our numerical experiments was around 85%, we initially explored the parameter space more freely by allowing (relatively) large jumps and gradually decreased  $\sigma_t^q$  as we made progress towards the feasible region; i.e. we adjusted the acceptance probability by allowing the tuning parameter to vary with  $t$ . According to Bounds (1987) “choosing an annealing schedule for practical purposes is still something of a black art” despite the simple principle underlying the choice of a suitable annealing schedule – the initial temperature should be low enough to “freeze” the system and should be increased towards its boiling point as the search progresses. For some parameters, nevertheless, one can certainly make a more informed decision. For multimodal problems, for example, using multiple (relatively) short chains would yield a better performance than using a single long chain.<sup>28</sup>

## 5 Conclusions

In this paper we propose a generalized risk budgeting (*GRB*) approach to portfolio construction. In comparison with the existing risk-based asset allocation techniques, our approach provides investors with more flexibility in that it allows investors to optimize a risk-return profile and to define risk budgets for possibly overlapping subsets of assets. Minimum variance, risk parity and risk budgeting strategies are therefore special cases of *GRB* strategies.

Although we show that the *GRB* problem can be formulated as a convex optimization problem in an important special case, the general *GRB* problem is a non-convex optimization problem. We introduce an SDP relaxation for bounding the optimal value of the *GRB* problem. When all assets have identical expected returns, our numerical results suggested that this SDP bound was quite tight, and could therefore be used to assess the quality of solutions produced by other approaches. Our main contribution in this paper is a simulation-based algorithm that combines augmented Lagrangian optimization ideas with MCMC methods. The goal of this algorithm is to compute a candidate solution in the neighborhood of the optimum, or a very good local optimal solution of the *GRB* problem. This candidate solution could then be used as the starting point for a standard non-linear optimization solver. In several numerical experiments our AL-MCMC algorithm was indeed successful in finding very good starting points.

We also note that our AL-MCMC approach is a general solution approach for solving non-convex optimization problems. The augmented Lagrangian algorithms is a very popular algorithm for computing local optimum solutions for non-convex problem. Combining this algorithm with the MCMC method opens up the possibility of converging to the *global* optimal solution, or at least providing a good starting-point for a non-linear optimization routine. In addition, this approach can be implemented very easily and is computationally fast.

We expect it to be of particular use for non-convex problems with small feasible regions where computing a good starting point is challenging. We intend to apply this approach to such problems in future research.

<sup>27</sup>For a discussion of this procedure in the context of MCMC, refer to Gelman, Carlin, Stern, Dunson, Vehtari, and Rubin (2004).

<sup>28</sup>Using multiple highly dispersed initial values to start several different chains is the most straightforward approach to solving a multimodal problem (Gelman and Rubin, 1992), and has a computational advantage for parallel processing machines.

## References

- A. Artzner, F. Delbaen, J.-M. Eber, and D. Heath. Coherent Measures of Risk. *Mathematical Finance*, 9(3):203 – 228, 1999.
- X. Bai, K. Scheinberg, and R. Tutuncu. Least-squares Approach to Risk Parity in Portfolio Selection. *Working Paper*, 2013.
- D. P. Bertsekas. *Constrained Optimization and Lagrange Method*. Athena Scientific, Belmont, MA, 1996.
- F. Black and R. Litterman. Global Portfolio Optimization. *Financial Analysts Journal*, 48:28 – 43, 1992.
- K. Boudt, P. Carl, and B. G. Peterson. Asset allocation with conditional value-at-risk budgets. *Journal of Risk*, 15(3):39–68, 2012.
- D. G. Bounds. New optimization methods from physics and biology. *Nature*, 329:215–218, 1987.
- S. Boyd and L. Vandenberghe. *Communications, Computation, Control and Signal Processing: A Tribute to Thomas Kailath*, chapter Semidefinite Programming Relaxations of Non-convex Problems in Control and Combinatorial Optimization, pages 279 – 288. Kluwer, 1997.
- B. Bruder and T. Roncalli. Managing Risk Exposures using the Risk Budgeting Approach. *SSRN*, 2012.
- V. DeMiguel, L. Garlappi, and R. Uppal. Optimal versus Naive Diversification: How Inefficient Is the 1/N Portfolio Strategy? *The Review of Financial Studies*, 22:1915 – 1953, 2009.
- A. Gelman and B. D. Rubin. Inferences from iterative simulation using multiple sequences (with discussion). *Statistical Science*, 7:457–511, 1992.
- A. Gelman, J. B. Carlin, H. S. Stern, D. B. Dunson, A. Vehtari, and D. B. Rubin. *Bayesian Data Analysis*. CRC Press, Boca Raton, FL, 2004.
- M. Grant and S. Boyd. Graph implementations for nonsmooth convex programs. In V. Blondel, S. Boyd, and H. Kimura, editors, *Recent Advances in Learning and Control*, Lecture Notes in Control and Information Sciences, pages 95–110. Springer-Verlag Limited, 2008. [http://stanford.edu/~boyd/graph\\_dcp.html](http://stanford.edu/~boyd/graph_dcp.html).
- M. Grant and S. Boyd. CVX: Matlab software for disciplined convex programming, version 2.1. <http://cvxr.com/cvx>, Mar. 2014.
- M. R. Hestenes. Multiplier and Gradient Methods. *Journal of Optimization Theory and Applications*, 4(5):303–320, 1969.
- S. Kirkpatrick, C. D. Gelatt, and M. P. Vecchi. Optimization by simulated annealing. *Science*, 220: 671–680, 1983.
- S. Maillard, T. Roncalli, and J. Teiletche. The Properties of Equally Weighted Risk Contribution Portfolios. *Journal of Portfolio Management*, 36(4):60 – 70, 2010.
- A. J. McNeil, R. Frey, and P. Embrechts. *Quantitative Risk Management: Concepts, Techniques, and Tools*. Princeton University Press, Princeton, NJ, 2005.



- Y. Nourani and B. Andresen. A comparison of simulated annealing cooling strategies. *Journal of Physics A: Mathematical and General*, 31:8373–8385, 1998.
- M. J. D. Powell. A Method for Nonlinear Constraints in Minimization Problems. In R. Fletcher, editor, *Optimization*. Academic Press, New York, NY, 1972.
- E. Qian. On the Financial Interpretation of Risk Contribution: Risk Budgets Do Add Up. *Journal of Investment Management*, 4(4):41 – 51, 2006.
- C. P. Robert and G. Casella. *Controlling Monte Carlo Variance*. Springer, New York, NY, 2004.
- G. O. Roberts, A. Gelman, and W. R. Gilks. Weak convergence and optimal scaling of random walk metropolis algorithms. Statistical Laboratory, Cambridge University, 1994.
- D. Ruppert. *Statistics and Data Analysis for Financial Engineering*. Springer, New York, NY, 2010.
- F. Spinu. An Algorithm for Computing Risk Parity Weights. *SSRN*, 2013.
- P. J. Van Laarhoven and E. H. L. Arts. *Simulated Annealing: Theory and Applications*. Reidel Publishers, Amsterdam, 1987.
- S. S. Zhu, D. Li, and X. L. Sun. Portfolio selection with marginal risk control. *Journal of Computational Finance*, 14(1):3–28, 2010.
- S. S. Zhu, a. X. T. Gui, X. L. Sun, and D. Li. Factor-risk-constrained mean-variance portfolio selection: Formulation and global optimization solution approach. *Journal of Risk*, 14(2):51–89, 2011/2012.

## A The Augmented Lagrangian Method

The augmented Lagrangian function was originally introduced by Hestenes (1969) and Powell (1972) to solve equality constrained optimization problems of the form:

$$\begin{aligned} \min_x \quad & F(x) \\ \text{subject to} \quad & h_k(x) = 0, \quad k = 1, \dots, s. \end{aligned} \quad (\text{A.1})$$

Instead of directly solving (A.1), the augmented Lagrangian method solves a sequence of problems of the form:

$$\min_x \quad \mathcal{L}_{c_t}(x, u_t) = F(x) + u_t' h(x) + \frac{1}{2} c_t \|h(x)\|_2^2 \quad (\text{A.2})$$

where the sequence of dual multipliers  $\{u_t\}$  are updated according to

$$u_{t+1} = u_t + c_t h(x_t),$$

where  $x_t \in \operatorname{argmin}_x \{\mathcal{L}_{c_t}(x, u_t)\}$ , and the sequence of penalty parameters  $\{c_t\}$  is a non-decreasing sequence. The augmented Lagrangian method was preceded by a pure penalty method where the multipliers  $u_t \equiv 0$ . Extensive empirical research showed that the pure penalty function method is often unreliable. The failure of this method is due to the fact that one needs to let  $c_t \rightarrow \infty$  in order to enforce feasibility; but computing the unconstrained minimization of  $\mathcal{L}_{c_t}(x, 0)$  becomes numerically difficult for large values of  $c_t$ .

On the other hand, the iterates computed by the augmented Lagrangian algorithm converge to feasible local optimal solution even for fixed penalty parameter  $c_t \equiv c$  (Bertsekas, 1996).

## B Numerical Results For Penalty-MCMC and Lagrangian-MCMC Methods

We now discuss numerical results for (i) the AL-MCMC method with  $u_t \equiv 0$  (the Penalty-MCMC method) and (ii) the AL-MCMC method with  $c_t \equiv 0$  (the Lagrangian-MCMC method). We compared the performance of these two methods with the full AL-MCMC method. In order to ensure that the comparison is fair, we use the same starting points for generating Markov chains across all three methods. Tables 6 and 7 report the results for the two scenarios: (i)  $\mu = \mu_0 \mathbf{1}$  and (ii)  $\mu \neq \mu_0 \mathbf{1}$ , respectively.

As can be seen from both tables, AL-MCMC-fmincon has a superior and more consistent performance when compared to Penalty-MCMC-fmincon and L-MCMC-fmincon. The range of objective values of the solutions computed by Penalty-MCMC-fmincon and Lagrangian-MCMC-fmincon are often very wide. For instance, when  $\mu \neq \mu_0 \mathbf{1}$ , in Set 1 Test Case 4, the solution ranges for Penalty-MCMC and Lagrangian-MCMC were  $[-102.97, -17.27]$  and  $[-102.47, -17.27]$ , respectively, while the range for AL-MCMC-fmincon is  $[-17.28, -17.27]$ . Also, in Set 1 Test Case 1 for the same scenario, the best solutions of Penalty-MCMC-fmincon and Lagrangian-MCMC-fmincon are 0.62 and 0.17 lower than the best solution of AL-MCMC-fmincon, respectively. Moreover, for higher dimensional problems, AL-MCMC-fmincon finds an optimal solution faster than the other methods. For example, in Set 2 Test Case 5, while AL-MCMC-fmincon took 40.66 seconds to yield an optimal solution, Penalty-MCMC-fmincon and L-MCMC-fmincon took 398.56 and 155.31 seconds, respectively. We can make similar observations when  $\mu = \mu_0 \mathbf{1}$ . We also remark that while failures in AL-MCMC-fmincon were caused by constraint violations that are less than  $10^{-4}$ , those in Penalty-MCMC-fmincon and L-MCMC-fmincon were caused by constraint violations that are

Table 6: Numerical Results for the Case of  $\mu = \mu_0 \mathbf{1}$ 

Table 6 presents numerical results for the three methods: AL-MCMC-fmincon, Penalty-MCMC-fmincon and L-MCMC-fmincon, when  $\mu = \mu_0 \mathbf{1}$ . The first and second columns contain the test set number and the name of the algorithm respectively. The third column reports the best solution obtained. The fourth column reports the range of the obtained solutions. The fifth column reports the number of solutions that failed to attain the same value as the best solution  $\min \tilde{F}(x^*)$ . The final column reports the amount of time taken to obtain the best solution.

Test Case 1: 7 Assets and 3 Subsets					
Set	Method	$\min \tilde{F}(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	33.09	[33.09, 34.16]	0	18.50
1	Penalty-MCMC-fmincon	33.16	[33.16, 64.81]	0	14.35
1	L-MCMC-fmincon	33.10	[33.10, 40.15]	0	14.88
2	AL-MCMC-fmincon	25.78	[25.78, 31.91]	0	16.90
2	Penalty-MCMC-fmincon	26.64	[26.64, 35.17]	2	9.28
2	L-MCMC-fmincon	29.79	[29.79, 31.70]	6	4.61

  

Test Case 2: 30 Assets and 5 Subsets					
Set	Method	$\min \tilde{F}(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	38.59	[38.59, 38.59]	0	2.54
1	Penalty-MCMC-fmincon	38.59	[38.59, 38.60]	0	2.91
1	L-MCMC-fmincon	38.59	[38.59, 38.60]	0	2.41
2	AL-MCMC fmincon	40.12	[40.12, 40.12]	0	0.81
2	Penalty-MCMC fmincon	40.12	[40.12, 40.12]	0	0.96
2	L-MCMC fmincon	40.12	[40.12, 40.12]	0	0.89

  

Test Case 3: 50 Assets and 5 Subsets					
Set	Method	$\min \tilde{F}(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	38.24	[38.24, 38.24]	0	2.68
1	Penalty-MCMC-fmincon	38.24	[38.24, 38.24]	1	2.40
1	L-MCMC-fmincon	38.24	[38.24, 38.24]	1	3.32
2	AL-MCMC fmincon	64.33	[64.33, 64.33]	0	1.90
2	Penalty-MCMC fmincon	64.33	[64.33, 64.33]	0	1.92
2	L-MCMC fmincon	64.33	[64.33, 64.33]	0	2.06

  

Test Case 4: 100 Assets and 10 Subsets					
Set	Method	$\min \tilde{F}(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	55.33	[55.33, 103.99]	0	7.98
1	Penalty-MCMC-fmincon	55.33	[55.33, 103.07]	0	10.02
1	L-MCMC-fmincon	55.33	[55.33, 55.86]	2	13.18
2	AL-MCMC fmincon	51.71	[51.71, 51.71]	0	7.51
2	Penalty-MCMC fmincon	51.71	[51.71, 51.71]	0	8.45
2	L-MCMC fmincon	51.71	[51.71, 51.71]	0	7.55

  

Test Case 5: 200 Assets and 10 Subsets					
Set	Method	$\min \tilde{F}(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	54.48	[54.48, 54.68]	5	67.89
1	Penalty-MCMC-fmincon	54.48	[54.48, 54.48]	0	76.37
1	L-MCMC-fmincon	54.47	[54.47, 54.51]	1	90.90
2	AL-MCMC-fmincon	55.02	[55.02, 55.23]	0	54.87
2	Penalty-MCMC fmincon	55.02	[55.02, 55.06]	1	55.94
2	L-MCMC fmincon	55.02	[55.02, 110.42]	0	64.93

greater than  $10^{-4}$ .

The results in this Appendix therefore demonstrate the advantage of incorporating the augmented Lagrangian method, as opposed to the penalty method or the Lagrangian multipliers method, into our MCMC algorithm.

Table 7: Numerical Results for the Case of  $\mu \neq \mu_0 \mathbf{1}$ 

Table 7 presents numerical results for the three methods: AL-MCMC-fmincon, Penalty-MCMC-fmincon and L-MCMC-fmincon, when  $\mu \neq \mu_0 \mathbf{1}$ . The first and second columns contain the test set number and the name of the algorithm respectively. The third column reports the best solution obtained. The fourth column reports the range of the obtained solutions. The fifth column reports the number of failures. The final column reports the amount of time taken to obtain the best solution.

Test Case 1: 7 Assets and 3 Subsets					
Set	Method	$\max F(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	63.89	[49.42, 63.89]	0	4.48
1	Penalty-MCMC-fmincon	63.27	[45.01, 63.27]	0	10.22
1	L-MCMC-fmincon	63.72	[54.06, 63.72]	0	21.47
2	AL-MCMC-fmincon	30.31	[30.31, 30.31]	0	0.58
2	Penalty-MCMC-fmincon	30.31	[30.31, 30.31]	1	0.94
2	L-MCMC-fmincon	30.31	[30.31, 30.31]	0	0.76
Test Case 2: 30 Assets and 5 Subsets					
Set	Method	$\max F(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	-25.33	[-25.33, -25.33]	0	2.37
1	Penalty-MCMC-fmincon	-25.33	[-25.33, -25.33]	0	3.44
1	L-MCMC-fmincon	-25.33	[-25.33, -25.33]	0	1.90
2	AL-MCMC fmincon	14.89	[14.89, 14.89]	0	0.68
2	Penalty-MCMC fmincon	14.89	[14.89, 14.89]	0	0.78
2	L-MCMC fmincon	14.89	[14.89, 14.89]	0	0.66
Test Case 3: 50 Assets and 5 Subsets					
Set	Method	$\max F(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	-0.30	[-0.30, -0.30]	0	2.52
1	Penalty-MCMC-fmincon	-0.30	[-0.30, -0.30]	1	2.34
1	L-MCMC-fmincon	-0.30	[-0.32, -0.30]	1	2.69
2	AL-MCMC fmincon	-20.72	[-20.72, -20.72]	0	2.00
2	Penalty-MCMC fmincon	-20.72	[-20.72, -20.72]	0	2.11
2	L-MCMC fmincon	-20.72	[-20.72, -20.72]	0	2.18
Test Case 4: 100 Assets and 10 Subsets					
Set	Method	$\max F(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	-17.27	[-17.28, -17.27]	0	7.43
1	Penalty-MCMC-fmincon	-17.27	[-102.97, -17.27]	0	7.38
1	L-MCMC-fmincon	-17.27	[-102.47, -17.27]	0	9.83
2	AL-MCMC fmincon	-0.46	[-0.46, -0.46]	0	7.61
2	Penalty-MCMC fmincon	-0.46	[-0.46, -0.46]	0	7.90
2	L-MCMC fmincon	-0.46	[-0.46, -0.46]	0	7.38
Test Case 5: 200 Assets and 10 Subsets					
Set	Method	$\max F(x^*)$	Solution range	No. of failures	$t$ (sec)
1	AL-MCMC-fmincon	8.41	[-172.35, 8.41]	4	47.55
1	Penalty-MCMC-fmincon	8.42	[8.41, 8.42]	0	50.93
1	L-MCMC-fmincon	8.42	[8.41, 8.42]	0	120.85
2	AL-MCMC-fmincon	7.52	[7.51, 7.52]	0	40.66
2	Penalty-MCMC fmincon	7.52	[7.50, 7.52]	0	398.56
2	L-MCMC fmincon	7.52	[-99.53, 7.52]	0	155.31